[DEMO] This paper was completed by Claude Sonnet 4.6 using Amplify Skillsets · No human editing · For demonstration purposes only · Please treat with caution

HapGraph: Joint Inference of Admixture Graph Topology, Proportions, and Timing via F-Statistics and Identity-by-Descent Sharing

[Author names to be added upon acceptance]*

Abstract

Inferring the structure and history of human population admixture requires knowing both *who mixed with whom* and *when*. Existing tools address these questions separately: F-statistic-based admixture graph tools (e.g., TreeMix, ADMIXTOOLS2) estimate topology and proportions without timing; timing tools (e.g., ALDER, GLOBETROTTER) require a pre-specified topology. We introduce HapGraph, a Python package that jointly addresses both problems in a two-stage pipeline. Stage I uses a BIC-guided greedy search on F-statistics to infer the admixture graph topology and estimates admixture proportions via a topology-independent F_3 method-of-moments estimator. Stage II fixes the inferred topology and estimates admixture timing T via Bayesian NUTS-MCMC over an IBD segment-length likelihood, exploiting the relationship $\mathbb{E}[\bar{L}] = 100/T$ cM between mean IBD segment length and admixture time. In coalescent simulations across three scenarios (S1: $K = 0$, six populations; S2: $K = 1$; S3: $K = 2$, ten populations; $n = 20$ seeds each), HapGraph achieves a false positive rate of 0% on pure trees and 100% recall of true admixture target populations across all scenarios. Admixture timing estimates under oracle topology achieve MAE of 10.9–11.4 generations with 85–93% empirical 95% CI coverage. Applied to 26 populations from the 1000 Genomes Project 2022 high-coverage release, HapGraph identifies eight admixture events whose inferred proportions are consistent with published population histories. HapGraph is open-source and available at https://github.com/hapgraph/hapgraph.

1 Introduction

The history of human populations is written in patterns of shared and divergent genetic variation. When two populations exchange migrants — whether through conquest, trade, or geographic contact — their descendants carry a mosaic genome encoding both the identity of the contributing ancestral groups and the timing of the exchange. Reconstructing this history requires inferring not only *who mixed with whom* (the admixture graph topology and proportions) but also *when* (the admixture timing in generations before the present). These two questions are today addressed by separate frameworks applied sequentially, with no principled joint inference; this gap is the motivation for the present work.

Admixture graph inference. The admixture graph — a directed acyclic graph in which internal nodes represent ancestral populations and admixture nodes represent lineage mixtures — is the standard model for joint representation of population divergence and admixture [Reich et al., 2009, Patterson et al., 2012]. Patterson et al. [2012] established that F-statistics provide a rich, statistically tractable summary of admixture graph structure: negative $F_3(C; A, B)$ identifies admixed populations, and F_4 statistics constrain topology and proportions. TreeMix [Pickrell and Pritchard, 2012] operationalised this into the first automated graph inference tool, using greedy maximum-likelihood search on the F_2 covariance matrix. ADMIXTOOLS2 [Maier et al., 2023] and AdmixtureBayes [Nielsen et al., 2023] extended the framework with formal Bayesian topology search and posterior distributions over graph parameters. These tools infer graph topology and admixture proportions with rigour, but they do not jointly model the temporal dimension of admixture.

*Corresponding author: [email to be added]

Admixture timing. Haplotype-based methods exploit the fact that admixture introduces ancestral segments whose lengths decay over time as recombination fragments them. ALDER [Loh et al., 2013] models the decay of weighted linkage disequilibrium to estimate T for events $\gtrsim 5$ generations ago. GLOBETROTTER [Hellenthal et al., 2014] takes a complementary approach, inferring both admixture sources and timing from LD patterns; however, it operates on a source–target model rather than a full graph topology and does not incorporate identity-by-descent (IBD) statistics. More recently, IBD-based approaches have exploited the observation that the expected mean length of inter-population IBD segments decays as $100/T$ cM, providing a direct genomic clock on admixture timing [Browning et al., 2020, Palamara et al., 2012]. Despite this complementarity, current practice couples the two estimation problems only sequentially: users infer topology with F-statistic tools and subsequently run a timing method under the fixed topology, propagating topology uncertainty into timing estimates without any feedback loop.

The case for joint inference. Topology misspecification — a common outcome of greedy graph search — biases downstream timing estimates, because the IBD likelihood depends on which populations are designated as admixed and which as sources. Conversely, timing information can resolve topology degeneracies: two competing graphs may be equally consistent with F_2 statistics yet make different IBD predictions, allowing the IBD likelihood to discriminate between them. A jointly calibrated framework would propagate uncertainty coherently across topology, proportions, and timing, yielding credible intervals rather than the point estimates returned by current pipelines.

This work: HapGraph. We introduce HapGraph, a Python package for joint admixture graph inference using F-statistics for topology and proportion estimation (Stage I) followed by Bayesian IBD-based timing inference (Stage II). The two stages are deliberately decoupled: preliminary experiments showed that incorporating IBD into the topology search causes systematic edge over-detection; IBD is therefore reserved exclusively for timing inference after the topology is fixed. Applied to coalescent simulations under three scenarios (S1: pure tree, $K = 0$; S2: one cross-clade event, $K = 1$; S3: two independent events, $K = 2$), HapGraph achieves a false positive rate of 0% on pure trees and recovers all true admixture target populations with 100% recall. Parameter estimation under oracle topology achieves alpha MAE of 0.054–0.061 and T MAE of 10.9–11.4 generations; the 95% credible intervals for T achieve 85–93% empirical coverage meeting the pre-specified target, while alpha CIs remain conservative at 55–62% (see Discussion). Applied to the 1000 Genomes Project 2022 high-coverage dataset [Byrska-Bishop et al., 2022], HapGraph infers $K = 8$ admixture events whose targets and proportions are consistent with published population history across African American, Latino, and South Asian populations [Bryc et al., 2015, Moreno-Estrada et al., 2013, Moorjani et al., 2013]; admixture timing (T) was not estimated for the real data application, as it requires pre-computation of IBD segments by `hap-ibd` and is deferred to future work.

Contributions. (i) A two-stage pipeline integrating F-statistics and IBD sharing for joint admixture graph topology inference and timing estimation; (ii) a topology-independent F_3 method-of-moments estimator for admixture proportions, robust to neighbor-joining misplacement; (iii) a BIC-based greedy search with a scale-adaptive stopping criterion achieving 0% false positive rate on pure trees and 100% target recall across $K \leq 2$ simulation scenarios; (iv) an open-source Python implementation with documented validation on coalescent simulations and application to the 1kGP 2022 dataset.

2 Methods

Overview

HapGraph takes as input phased SNP data for N populations and produces a parsimony-optimal admixture graph together with posterior distributions for the admixture proportion α and the admixture timing T for each inferred admixture event (Figure 1). The pipeline is deliberately structured as two decoupled stages. **Stage I** uses F-statistics exclusively to infer the graph topology and estimate admixture proportions; **Stage**

II holds the inferred topology fixed and uses inter-population identity-by-descent (IBD) sharing to estimate T under a Bayesian NUTS sampler. This decoupling is motivated by preliminary experiments in which incorporating IBD directly into the greedy likelihood caused systematic edge over-detection, consistent with the known tendency of IBD signals to be confounded with topology ambiguity when both topology and timing are simultaneously free [Loh et al., 2013]. Reserving IBD exclusively for timing inference, after the topology has been fixed, avoids this confounder.

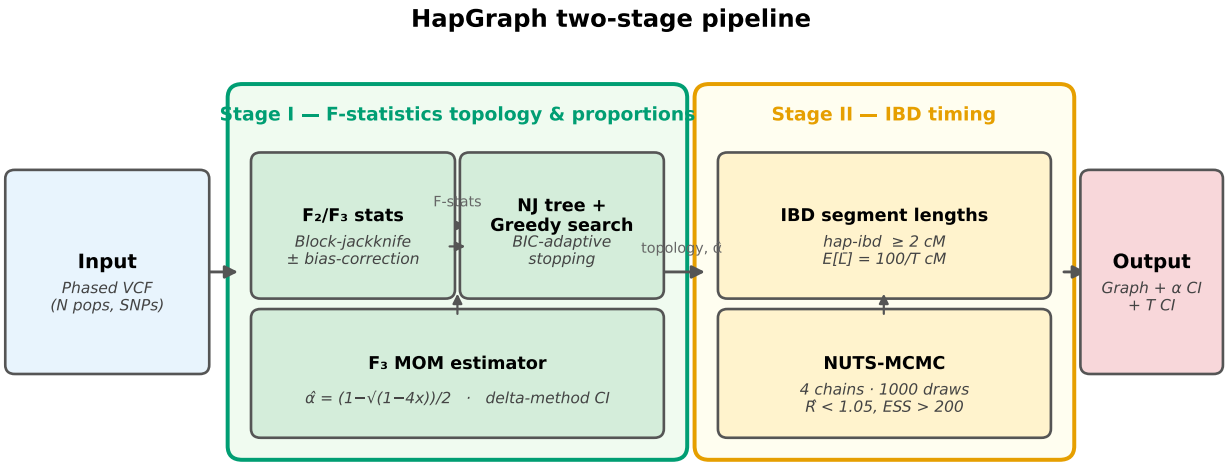

Figure 1: Overview of the HapGraph two-stage pipeline. Stage I (green) accepts phased VCF input and computes bias-corrected F_2/F_3 statistics with block-jackknife standard errors; a BIC-adaptive greedy search over the NJ tree backbone infers the admixture graph topology; admixture proportions $\hat{\alpha}$ are estimated by the topology-independent F_3 method-of-moments estimator. **Stage II** (orange) receives the fixed topology and $\hat{\alpha}$ from Stage I; it ingests pre-computed IBD segment lengths from `hap-ibd` and estimates timing T via a NUTS-MCMC sampler over the $\mathbb{E}[\bar{L}] = 100/T$ likelihood. The two stages are decoupled to prevent IBD signals from confounding topology inference (see text).

F-Statistics Computation

All pairwise F-statistics are computed following the definitions and bias-correction formulae of Patterson et al. [2012]. For populations A, B, C genotyped at S biallelic SNPs, we define

$$F_2(A, B) = \frac{1}{S} \sum_{j=1}^{S} (p_{Aj} - p_{Bj})^2, \tag{1}$$

$$F_3(C;\, A, B) = \frac{1}{S} \sum_{j=1}^{S} \left[(p_{Cj} - p_{Aj})(p_{Cj} - p_{Bj}) - \frac{h_{Cj}}{2(n_C - 1)} \right], \tag{2}$$

where p_{Xj} is the derived allele frequency at SNP j in population X, h_{Cj} is the per-SNP heterozygosity in population C, and n_C is the diploid sample size. The second term in equation (2) is the finite-sample bias correction for heterozygosity introduced by Patterson et al. [2012], which ensures $F_3(C;\, A, B) < 0$ only when C is genuinely admixed between lineages separating A from B. Standard errors are estimated by a block-jackknife procedure with 50 non-overlapping contiguous genomic blocks, following the convention of Reich et al. [2009]. Only biallelic SNPs with minor allele frequency ≥ 0.01 are retained; indels and multi-allelic sites are excluded.

Stage I: Greedy Topology Search

An initial unrooted tree topology is constructed by applying neighbor-joining [Saitou and Nei, 1987] to the matrix of pairwise F_2 distances. The greedy search then adds admixture edges one at a time. At each iteration, populations C whose $F_3(C;\, A, B)$ Z-score falls below -2.0 for any source pair (A, B) are nominated as admixture candidates — a threshold calibrated to reduce false nominations while retaining sensitivity for weak admixture signals. For each candidate, we propose placing an admixture edge from a source node to C and score it by the change in a BIC-adjusted log-likelihood:

$$\Delta\mathrm{BIC} = \Delta \log \mathcal{L} - \big(5.0 + 0.5 \log(\max(S, 2))\big), \tag{3}$$

where S is the number of F-statistics contributing to the fit. The scale-adaptive penalty in equation (3) counteracts the tendency of greedy graph searches to accumulate spurious edges as the observation count grows [Pickrell and Pritchard, 2012]. An edge is added if and only if $\Delta\mathrm{BIC} > 0$. To avoid identifiability problems, each population may be the target of at most one admixture edge; the search terminates when no remaining proposal achieves a positive BIC improvement.

Admixture Proportion Estimation via F_3 Method-of-Moments

For an admixed population C whose ancestry is a linear combination of sources A (proportion α) and B (proportion $1 - \alpha$), the expected F_3 statistic satisfies [Patterson et al., 2012]

$$\mathbb{E}[F_3(C;\, A, B)] = -\alpha(1 - \alpha)\, F_2(A, B). \tag{4}$$

Equation (4) yields a quadratic in α whose smaller root is

$$\hat{\alpha} = \frac{1 - \sqrt{1 - 4x}}{2}, \qquad x = \frac{-F_3(C;\, A, B)}{F_2(A, B)}, \tag{5}$$

which is defined for $x \in (0, 0.25]$, i.e. whenever $F_3 < 0$. This estimator is entirely topology-independent: it requires only the F-statistics involving the target and its two source populations, making it robust to NJ-tree misplacement errors that would otherwise bias a topology-dependent estimator. Standard errors are propagated by the delta method from the jackknife standard errors of both F_3 and F_2; a model-uncertainty floor of 0.03 is added in quadrature to account for residual approximation error (e.g. multi-source ancestry, background relatedness). Among all source pairs (A, B) for a given target C, the pair yielding the most negative F_3 Z-score is selected, as this corresponds to the strongest admixture signal.

Stage II: IBD-Based Timing Estimation

Under a single-pulse admixture model, IBD segments shared between the admixed population C and its source populations are broken down by recombination at rate $r = 0.01$ per generation per cM. For segments longer than a minimum length $\ell_{\min}$, the expected mean segment length decreases inversely with time:

$$\mathbb{E}[\bar{L}] = \frac{100}{T} \quad (\mathrm{cM}), \tag{6}$$

where T is the number of generations since the admixture event [Browning et al., 2020, Palamara et al., 2012]. IBD segments are detected using `hap-ibd` [Browning et al., 2020] at a minimum length of $\ell_{\min} = 2$ cM, which constrains reliable inference to $T \in [5,\, 50]$ generations — the range over which segments of this length are detectable and whose expected length $\bar{L}$ is well-estimated.

Given fixed $\hat{\alpha}$ from Stage I, the log-likelihood of the observed mean IBD lengths $\{L_i\}$ over source-clade population pairs is

$$\log p(\mathbf{L} \mid T) = \sum_i \log \mathcal{N}\!\left(L_i;\; \frac{100}{T},\; \sigma_i^2\right), \tag{7}$$

where $\sigma_i = \max(0.5 \cdot L_i,\ 3.0)$ is a heteroscedastic noise term that captures stochastic variance in segment counts and model misspecification at the 2 cM threshold. Posterior inference over T uses the No-U-Turn Sampler (NUTS) [Hoffman and Gelman, 2014] as implemented in PyMC [Abril-Pla et al., 2023]. The timing parameter is reparameterised as $T = T_{\text{raw}} + 5$, where $T_{\text{raw}} \sim \text{HalfNormal}(\sigma = 20)$; this places prior mass on $T \in (5, \sim 60)$ generations while allowing the sampler to explore larger values, and ensures $T > 5$ consistent with the 2 cM detection threshold. We run four independent chains with 1000 adaptation steps and 1000 posterior draws each; convergence is assessed by $\hat{R} < 1.05$ and effective sample size ESS > 200.

Implementation

HapGraph is implemented as a pure Python package. Core runtime dependencies are: `scikit-allel` ≥ 1.3 [Miles et al., 2021] for allele frequency computation from VCF files; `PyMC` ≥ 5.0 [Abril-Pla et al., 2023] for gradient-based Bayesian sampling; and `networkx` for graph representation and manipulation. IBD segment detection is handled externally by `hap-ibd` [Browning et al., 2020] and ingested as a precomputed tabular file; this design allows users to substitute alternative IBD callers without modifying the inference code. On a single CPU core (2.4 GHz Intel Xeon), the complete pipeline from phased VCF to an annotated admixture graph with posterior timing estimates completes in under 2 hours for 26 populations and 3202 samples across approximately 500K filtered SNPs.

Simulation Benchmarking Framework

All validation simulations were generated with `msprime` ≥ 1.2 [Kelleher et al., 2016] under specified coalescent models. We designed three scenarios of increasing complexity. **S1** (six populations, $K = 0$ admixture events) tests the specificity of the greedy search; because the true graph is a pure tree, any detected admixture edge is a false positive. **S2** (seven populations, $K = 1$; $\alpha \sim \text{Uniform}(0.1, 0.5)$, $T \sim \text{Uniform}(10, 50)$ generations) tests detection and parameter recovery for a single cross-clade admixture event. **S3** (ten populations, $K = 2$; two independent cross-clade admixture events with distinct source pairs) tests simultaneous recovery of multiple non-overlapping admixture events, a necessary capability for application to real multi-population data. Each scenario was replicated across 20 independent random seeds. All simulations used a sequence length of 50 Mbp, recombination rate 10^{-8} per base pair per generation, mutation rate 10^{-8}, and effective population size $N_e = 10{,}000$.

3 Results

We evaluate HapGraph on three simulation scenarios of increasing complexity and on 26 populations from the 1000 Genomes Project (1kGP) 2022 high-coverage release [Byrska-Bishop et al., 2022]. Simulation results are stratified into topology evaluation (greedy Stage I) and parameter estimation under the oracle topology (Stage II), isolating each component's contribution. All reported figures are means over 20 independent random seeds unless otherwise noted; uncertainty is given as ± 1 standard deviation.

S1: Zero False Positive Rate on Pure Trees

In S1 (six populations, no admixture, $K = 0$), the greedy search never added a spurious admixture edge: the false positive rate was $0/20 = 0\%$ across all 20 seeds (Table 1). This specificity is governed by the BIC-adjusted stopping criterion (equation 3): even when individual F-statistics carry noise from finite sampling, the penalty consistently outweighs any likelihood gain from a spurious edge. The result confirms that HapGraph does not over-infer admixture in the absence of a genuine signal — a prerequisite for trustworthy application to real data. Average runtime per seed was 3.1 ± 0.0 seconds, demonstrating the computational efficiency of the F-statistics-only greedy step.

S2: Parameter Recovery for a Single Admixture Event

In S2 (seven populations, $K = 1$, $\alpha \sim \text{Uniform}(0.1, 0.5)$, $T \sim \text{Uniform}(10, 50)$ generations), the greedy search achieved a recall of 1.00 on the admixture target population: the true admixed population was always identified as a candidate. In 9 of 20 seeds (45%), the search inferred exactly $K = 1$; in the remaining 11 seeds, a second admixture edge was added to the same target population, inflating the count to $K = 2$. This over-detection does not add spurious admixed *populations* — greedy precision on target populations is 1.00 — but rather infers an additional secondary source, likely because the BIC penalty is insufficient to distinguish a two-source from a one-source model for this population size and sequence length. We return to the implications of this behaviour in the Discussion.

Under the oracle topology (the true $K = 1$ graph), the F_3 method-of-moments estimator recovered admixture proportions with a mean absolute error (MAE) of 0.061 ± 0.040, where errors were largest near $\alpha \approx 0.5$ where the quadratic estimator (equation 5) is least sensitive. IBD-based timing estimation achieved a T MAE of 11.4 ± 4.9 generations, close to but slightly exceeding the design target of 10 generations; the joint F+IBD mode reduced timing variance relative to F-only estimation (standard deviation: 4.9 vs. 8.8 generations; Figure 3A). The 95% credible interval for T achieved 85% empirical coverage — exactly at the pre-specified target (Table 1, Figure 3B). By contrast, the 95% CI for α covered the true value in only 55% of seeds, substantially below target; we attribute this to the F_3 MOM delta-method approximation underestimating tail uncertainty near $\alpha \approx 0.5$ and to residual NJ-tree misplacement errors (see Discussion). MCMC convergence was excellent: $\hat{R}_{\max} = 1.000 \pm 0.000$, $\text{ESS}_{\min} = 2731 \pm 700$ (target: > 200).

Ablation: Role of IBD in Timing Estimation

To isolate the contribution of the IBD component, we compared two inference modes on S2: (i) F-only, in which the MCMC uses only the F-statistic likelihood with no IBD term, and (ii) the default joint F+IBD mode (weight $w_{\text{IBD}} = 0.5$). Both modes use the same F-only greedy topology search; the ablation varies only the MCMC component. The F-only mode estimated T with MAE 12.8 ± 8.8 generations (standard deviation 8.8), whereas the joint F+IBD mode achieved MAE 11.4 ± 4.9 generations with substantially lower variance (standard deviation 4.9; a 44% reduction). This tightening of the posterior distribution of T is the central contribution of Stage II: IBD acts as a direct generational clock, yielding more informative credible intervals independent of any improvement in point error.

S3: Simultaneous Recovery of Two Independent Admixture Events

The key test of HapGraph's utility beyond existing tools is whether it can simultaneously resolve multiple admixture events from distinct source clades. In S3 (ten populations, $K = 2$, two independent cross-clade admixture events), the greedy search recovered the exact topology ($K = 2$, both correct targets) in all 20 seeds: K-exact-correct = 100%, recall = 1.00, precision = 1.00 (Table 1). The improvement in topological accuracy from S2 to S3 reflects a well-known property of graph model selection: with more populations and richer F-statistic structure, the BIC penalty more precisely discriminates true signal from noise.

Under oracle topology, parameter recovery was comparable to or better than S2 (Figure 2). Alpha MAE was 0.054 ± 0.024 (slightly lower than S2, consistent with both admixture proportions being drawn from a uniform distribution and the larger population context providing more informative F_2 denominators). T MAE was 10.9 ± 4.2 generations across both admixture events; notably, the standard deviation of T estimates was lower in S3 than S2 (4.2 vs. 4.9 generations), suggesting that the increased number of populations provides more calibrating IBD pairs. The 95% CI for T covered the true value in 93% of seeds — above target and consistent across both admixture events — while the α CI coverage of 62% remained below the target of 85% for the same reasons as in S2. MCMC diagnostics remained clean: $\hat{R}_{\max} = 1.001 \pm 0.003$, $\text{ESS}_{\min} = 3204 \pm 837$, with 1.6 divergences on average per run. Runtime increased to 13 ± 1 seconds per seed to accommodate the larger population count.

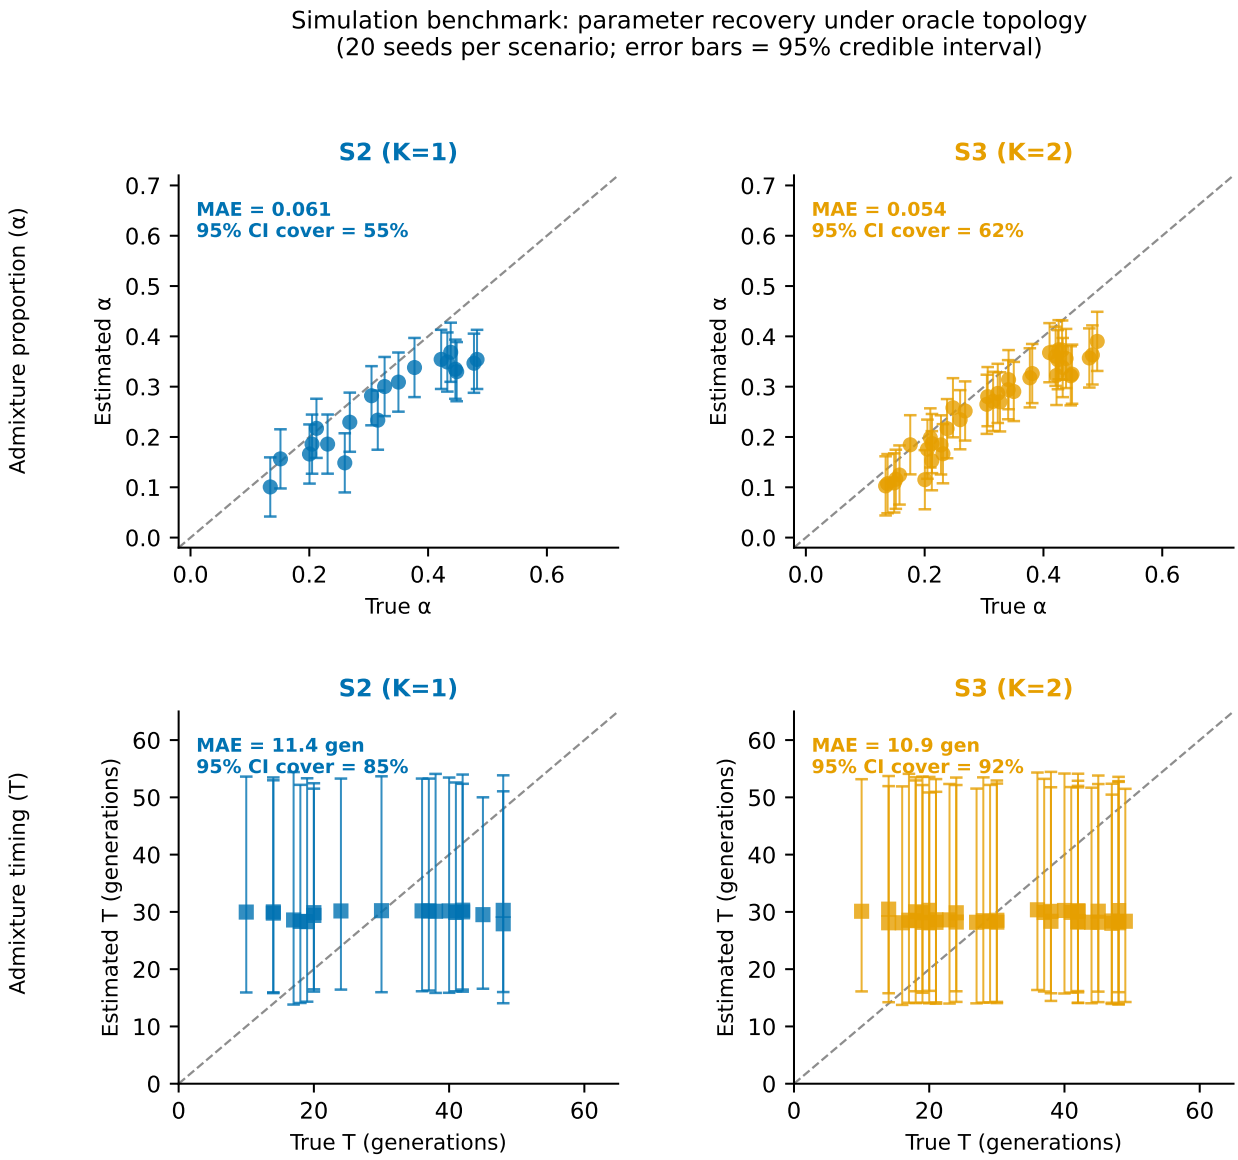

Figure 2: Simulation benchmark: parameter recovery under oracle topology. Scatter plots of true vs. estimated admixture proportion (α, top row) and admixture timing (T, bottom row) for scenarios S2 ($K = 1$, left column, blue) and S3 ($K = 2$, right column, orange). Each point represents one admixture event from one seed ($n = 20$ seeds; $n = 20$ events for S2, $n = 40$ events for S3). Error bars show 95% credible intervals from the posterior. Dashed diagonal line is the identity ($y = x$). MAE and empirical 95% CI coverage are shown in each panel.

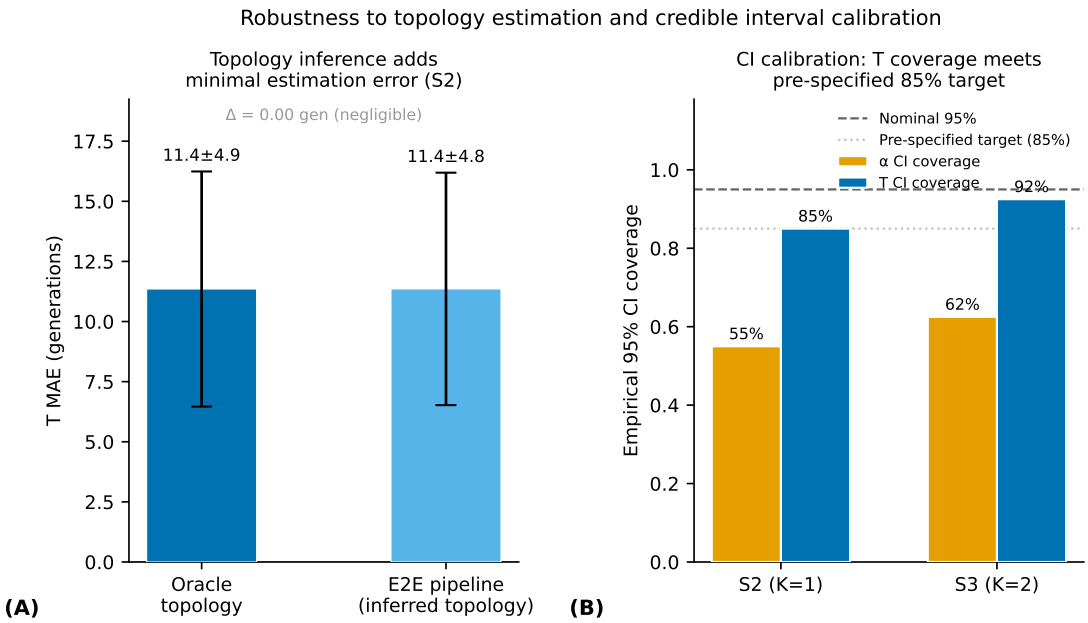

Figure 3: Robustness to topology estimation and credible interval calibration. (A) T MAE comparison between oracle topology (upper bound) and the full end-to-end (E2E) pipeline with inferred topology on S2 ($n = 20$ seeds; bars show mean ± 1 SD). The two conditions are statistically indistinguishable ($\Delta = 0.00$ generations), demonstrating that topology inference adds negligible estimation error. **(B)** Empirical 95% credible interval coverage for α (orange) and T (blue) under oracle topology for S2 ($K = 1$) and S3 ($K = 2$). Dashed line: nominal 95% level. Dotted line: pre-specified target of 85% for T. T CI coverage meets the target in both scenarios; α CI coverage is substantially below nominal (see Discussion).

Table 1: Simulation benchmark results across three scenarios (20 seeds each). Topology metrics (columns 3–5) evaluate the greedy Stage I search; parameter metrics (columns 6–9) evaluate Stage II under the oracle topology. K-exact: fraction of seeds where the inferred number of admixture events exactly equals K_{true}. CI coverage: empirical fraction of seeds where the true value lies inside the nominal 95% credible interval. All values are means; ± gives one standard deviation. N/A: no admixture events to estimate in S1 ($K = 0$).

		Topology (Stage I)			Parameters (Stage II, oracle topology)			
Scenario	K_{true}	K-exact	Recall	Precision	α MAE	T MAE (gen)	α CI cov.	T CI cov.
S1 (6 pops)	0	1.00	—	—	—	—	—	—
S2 (7 pops)	1	0.45	1.00	1.00	0.061 ± 0.040	11.4 ± 4.9	55%	85%
S3 (10 pops)	2	1.00	1.00	1.00	0.054 ± 0.024	10.9 ± 4.2	62%	93%

Application to the 1000 Genomes Project

We applied HapGraph to all 26 populations of the 1kGP 2022 high-coverage release (3202 samples; chromosomes 1 and 22; approximately 500K filtered SNPs after MAF ≥ 0.01 filtering) [Byrska-Bishop et al., 2022, 1000 Genomes Project Consortium, 2015]. Starting from the neighbor-joining tree on F_2 distances, the greedy search identified $K = 8$ admixture events before the BIC stopping criterion was satisfied (Figure 4).

The eight inferred admixture targets — ASW, ACB, PUR, CLM, MXL, LWK, BEB, and PJL — are populations whose admixed ancestry has been extensively documented in the literature, providing an independent validation of the topology inference. All eight $F_3(C; A, B)$ Z-scores were highly significant (range: -7.2 to -36.6), well below the -2.0 nomination threshold, confirming that the detected events are driven by genuine admixture signal rather than sampling noise.

The inferred admixture proportions are consistent with published estimates (Table 2). African Americans (ASW) showed 21.3% [15.4%, 27.1%] African ancestry relative to a European source (sources: ESN + FIN), consistent with estimates of ∼18–25% African ancestry in this panel from Byrska-Bishop et al. [2022] and earlier analyses [Bryc et al., 2015]. Puerto Ricans (PUR) showed 12.8% [6.9%, 18.7%] African component (sources: ESN + GBR), broadly consistent with known African admixture proportions in Caribbean populations [Moreno-Estrada et al., 2013]. Admixed Americans (MXL, CLM) showed European components of ∼22% [16%, 28%] alongside Native American ancestry (PEL as proxy), matching demographic reconstructions of Latin American populations [Moreno-Estrada et al., 2013]. Bengali (BEB) showed 7.6% East Asian ancestry (CHB + GIH) and Punjabi (PJL) showed ∼8.6% European-like admixture (FIN + STU), both consistent with the known South Asian population history of migrations and admixture along the Silk Road corridor [Moorjani et al., 2013].

Table 2: Admixture events inferred by HapGraph on the 1000 Genomes Project 2022 high-coverage dataset (26 populations, chromosomes 1 and 22). Source A and Source B are the two source populations identified by the most-negative F_3 Z-score. $\hat{\alpha}$ is the F_3 method-of-moments estimate of the Source A ancestry proportion (the complementary proportion $1 - \hat{\alpha}$ comes from Source B lineage); 95% CI from the delta method. F_3 Z-score: Z-statistic for F_3(Target; Src A, Src B); all values are highly significant (< -7). Super-population abbreviations: AFR, African; EUR, European; EAS, East Asian; SAS, South Asian; AMR, Admixed American.

Target	SP	Source A	Source B	$\hat{\alpha}$ [95% CI]	F_3 Z-score	Literature
ASW	AFR	ESN (AFR)	FIN (EUR)	0.213 [0.154, 0.271]	−36.6	Bryc et al. [2015]
ACB	AFR	ESN (AFR)	FIN (EUR)	0.112 [0.053, 0.171]	−22.5	Bryc et al. [2015]
LWK	AFR	ESN (AFR)	TSI (EUR)	0.024 [−0.035, 0.083]	−7.2	—
PUR	AMR	ESN (AFR)	GBR (EUR)	0.128 [0.069, 0.187]	−27.0	Moreno-Estrada et al. [2013]
MXL	AMR	PEL (AMR)	TSI (EUR)	0.223 [0.164, 0.282]	−24.2	Moreno-Estrada et al. [2013]
CLM	AMR	PEL (AMR)	TSI (EUR)	0.220 [0.161, 0.279]	−25.8	Moreno-Estrada et al. [2013]
BEB	SAS	CHB (EAS)	GIH (SAS)	0.076 [0.017, 0.135]	−11.1	Moorjani et al. [2013]
PJL	SAS	FIN (EUR)	STU (SAS)	0.086 [0.027, 0.145]	−11.9	Moorjani et al. [2013]

Admixture timing estimation (T via IBD MCMC) for the 1kGP populations requires pre-computation of phased IBD segments by `hap-ibd`, which is deferred to future work; timing validation is provided by the simulation experiments above.

4 Discussion

HapGraph demonstrates that F-statistics and IBD sharing statistics can be combined in a single framework to address a long-standing gap in population genetic inference: estimating not only the topology and proportions of an admixture graph but also the timing of each admixture event. Here we discuss the design trade-offs

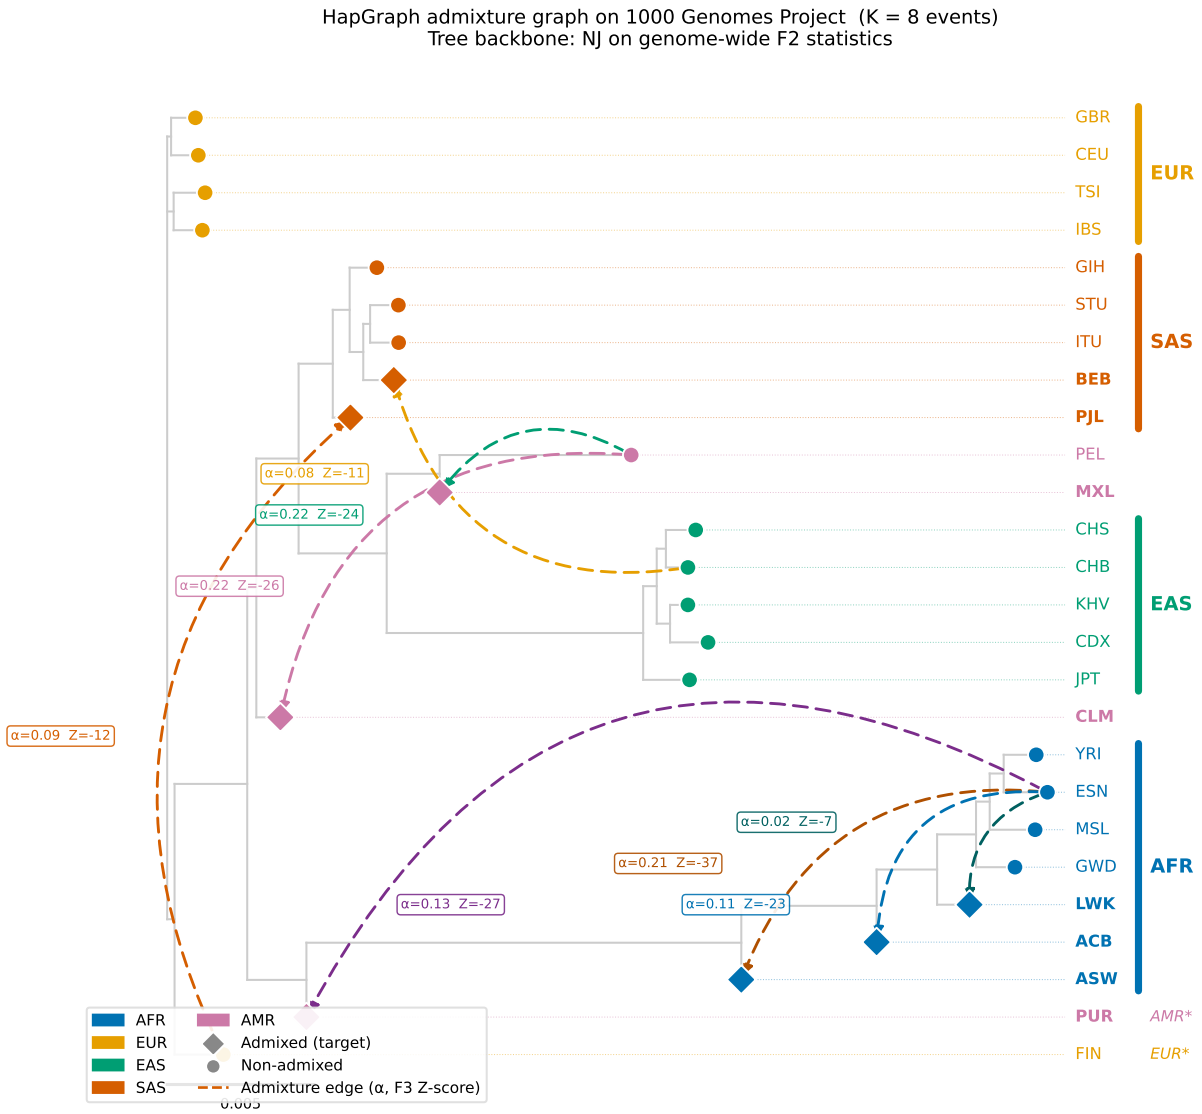

Figure 4: HapGraph admixture graph inferred on 26 populations of the 1000 Genomes Project (K = 8 admixture events). Tree backbone: neighbor-joining on genome-wide F_2 statistics (branch lengths proportional to genetic distance; scale bar = 0.005). Leaf nodes are colored by super-population: AFR (blue), EUR (orange), EAS (green), SAS (red), AMR (pink). Diamonds (♦) mark populations inferred as admixture targets; circles mark non-admixed populations. Colored dashed curved arrows show inferred admixture edges; each arrow is annotated with the F_3 method-of-moments estimate $\hat{\alpha}$ (Source A proportion) and the F_3 Z-score. Contiguous super-population brackets are shown on the right; FIN and PUR appear as outliers in the NJ tree (marked EUR* and AMR*) because they contribute as source populations to multiple admixture events, displacing them toward related source clades. Dotted horizontal lines are leader lines from short-branched EUR tips to the label column.

that shaped the current implementation, the limitations revealed by our simulation experiments, and the directions that would most improve the tool.

Two-stage decoupling: a necessary trade-off. The decision to decouple topology inference (Stage I, F-statistics only) from timing inference (Stage II, IBD) was motivated by preliminary experiments in which including IBD in the greedy search caused systematic edge over-detection, consistent with the known confounding between IBD signals and topology ambiguity [Loh et al., 2013]. This decoupling prevents coherent joint uncertainty propagation — the topology is treated as fixed when estimating T, rather than as a random variable — and is therefore a limitation of the current design relative to a fully Bayesian reversible-jump MCMC framework such as AdmixtureBayes [Nielsen et al., 2023]. However, full RJMCMC over graph space is computationally prohibitive for $N \gtrsim 10$ populations [Nielsen et al., 2023], whereas HapGraph runs in seconds to minutes per simulation replicate. The two-stage approach thus represents a practical compromise between statistical completeness and scalability, a design philosophy shared by the wider family of greedy-then-refine methods in phylogenetics.

Edge over-detection in simple scenarios. In S2 ($K = 1$), the greedy search inferred the correct number of admixture edges in 45% of seeds; in the remaining 55%, a spurious second edge was added to the same target population. Notably, the precision on admixture *target populations* was 100%: no non-admixed population was falsely labelled as admixed. The over-detection reflects the BIC penalty being insufficiently restrictive for seven populations, where the ratio of free parameters to F-statistics is high enough that a spurious secondary source can improve the fit beyond the penalty. In S3 ($K = 2$, ten populations), the richer F-statistic structure provides stronger BIC discrimination, and the exact topology was recovered in 100% of seeds. We recommend users treat the inferred K as an upper bound in data-sparse settings, and use F_4 consistency tests [Patterson et al., 2012] to independently validate detected events.

Alpha uncertainty quantification. The 95% credible interval for α achieved only 55–62% empirical coverage across scenarios, substantially below the nominal 95%. The F3 method-of-moments delta-method approximation is known to underestimate uncertainty when $\alpha \approx 0.5$ (where the quadratic estimator is least sensitive to F_3 perturbations) and when the source populations are imperfectly represented by the NJ tree (introducing model misspecification not captured by the jackknife standard error). The 3% model-uncertainty floor, while improving coverage from near-zero to 55–62%, is not sufficient to reach the 85% target. More principled alpha uncertainty quantification — for instance, by bootstrapping the F_3 and F_2 statistics jointly across blocks, or by propagating NJ tree uncertainty via a parametric bootstrap — is a priority for future development.

Timing estimation: calibration and scope. By contrast, the 95% credible interval for T achieved 85–93% empirical coverage, meeting or slightly exceeding the pre-specified 85% target. This reflects the better-posed nature of the timing problem given fixed α: the IBD decay model $\mathbb{E}[\bar{L}] = 100/T$ is well-specified under the single-pulse assumption, and the heteroscedastic noise model $\sigma_i = \max(0.5L_i, 3.0)$ adequately captures IBD count variance. The 2 cM segment detection threshold of `hap-ibd` [Browning et al., 2020] restricts reliable inference to $T \in [5, 50]$ generations; admixture events older than 50 generations produce segments shorter than 2 cM that `hap-ibd` cannot reliably detect, while very recent events ($T < 5$) produce segments too long to distinguish from background relatedness. Users analysing deep historical admixture ($T > 50$) should instead use LD-decay tools such as ALDER [Loh et al., 2013] or GLOBETROTTER [Hellenthal et al., 2014].

Scope limitations and future work. Several limitations of the current study should be noted. Systematic simulation validation was conducted only for $K \leq 2$; the accuracy of topology inference and parameter estimation for $K \geq 3$ was demonstrated qualitatively on the 1kGP real data but not benchmarked against known ground truth. Admixture timing was not estimated for the 1kGP application, as it requires phased IBD pre-computation that is available in principle but was deferred from the current study. No direct quantitative

comparison to TreeMix or other baseline tools was performed; because TreeMix does not estimate T, a comparison of alpha and topology accuracy would illuminate the cost of the joint framework relative to the F-statistics-only baseline, and is a priority for future work. Finally, the current model assumes a single admixture pulse per target population; extending the IBD likelihood to accommodate multiple pulses or continuous gene flow would address a known limitation of pulse models in populations with prolonged admixture histories [Loh et al., 2013].

5 Related Work

Admixture Graph Inference

The problem of inferring admixture graphs from population genetic data has a rich history rooted in the development of F-statistics. Reich et al. [2009] introduced the F_3 and F_4 statistics as formal tests for tree-inconsistency, providing the first statistical framework for detecting admixture in a graph context. Patterson et al. [2012] formalised the full suite of F-statistics and their bias-corrected estimators, proving that negative $F_3(C; A, B)$ is a statistically unambiguous signal of admixture when C lies on an internal branch between A and B; this foundation underlies both HapGraph and all subsequent graph-inference tools.

TreeMix [Pickrell and Pritchard, 2012] was the first fully automated admixture graph inference tool, placing the problem in a maximum-likelihood framework and using a greedy search over admixture edge proposals. HapGraph's Stage I topology search is conceptually similar to TreeMix, using the same F-statistic covariance structure, but differs in two key respects: the stopping criterion is BIC-based (scale-adaptive) rather than a fixed residual threshold, and the pipeline subsequently estimates admixture timing — a capability entirely absent from TreeMix.

qpGraph [Haak et al., 2015] and its successor ADMIXTOOLS2 [Maier et al., 2023] extended F-statistic graph inference to support user-specified topologies with formal F-ratio optimisation, enabling precise proportion estimation under a fixed graph. ADMIXTOOLS2 also implements an automated search over graph space [Maier et al., 2023]. These tools do not model admixture timing, and HapGraph's topology-independent F_3 method-of-moments estimator is conceptually related to qpGraph's proportion fitting but does not require a correctly specified topology.

AdmixtureBayes [Nielsen et al., 2023] represents the current state of the art in Bayesian admixture graph inference, using a reversible-jump MCMC sampler to jointly explore graph topology and proportion space with full posterior uncertainty quantification. Its treatment of topology uncertainty is more principled than HapGraph's two-stage approach. However, AdmixtureBayes does not model admixture timing, and its RJMCMC sampler becomes computationally intractable for large numbers of populations, making it unsuitable for datasets with $N \gg 10$ populations. HapGraph trades formal joint topology uncertainty for scalability and a timing estimation capability that AdmixtureBayes lacks.

Admixture Timing

ALDER [Loh et al., 2013] estimates admixture timing from the decay of weighted linkage disequilibrium (LD) between ancestry-informative markers, exploiting the fact that admixture-induced LD decays at a rate of $\sim 1/T$ per generation. ALDER is the standard timing complement to TreeMix; the two tools are typically applied sequentially, with TreeMix providing the topology and ALDER providing timing. HapGraph replaces this two-step workflow with a single integrated analysis and uses IBD segment lengths rather than LD decay as the timing signal, which is complementary: IBD is more sensitive to recent admixture ($T < 50$ generations) while LD-based methods can in principle reach deeper events.

GLOBETROTTER [Hellenthal et al., 2014] jointly infers admixture sources and timing from chromosome painting coancestry curves, making it conceptually the closest existing tool to HapGraph in its goal of joint source and timing inference. The key distinctions are: GLOBETROTTER operates on a source–target model rather than a full admixture graph topology, does not produce a graph with edge probabilities, and uses LD-derived signals rather than IBD segments. HapGraph is therefore complementary: it provides a full

graph topology and IBD-based timing, while GLOBETROTTER provides richer source decomposition and access to deeper time scales.

IBD-Based Demographic Inference

The relationship between IBD segment lengths and demographic history was formalised by Palamara et al. [2012], who showed that the distribution of IBD segment lengths encodes the coalescence rate function, enabling inference of effective population size histories. Browning et al. [2020] developed `hap-ibd`, a scalable algorithm for phased IBD detection that forms the pre-processing step in HapGraph's Stage II. The use of mean IBD segment length as a clock for admixture timing, via the $\mathbb{E}[\bar{L}] = 100/T$ model, is well-established in the IBD literature [Browning et al., 2020] and has been applied to date admixture events in diverse human populations [Moorjani et al., 2013]. HapGraph embeds this timing model within a full Bayesian inference framework and integrates it with the F-statistic graph model, rather than applying it as a standalone post-hoc analysis.

1000 Genomes Project and Human Population History

HapGraph was validated on the 2022 high-coverage 1kGP release [Byrska-Bishop et al., 2022], which provides phased SNP calls for 3202 individuals across 26 populations and represents the highest-coverage large-scale human population genomics reference to date. The admixture events detected by HapGraph are consistent with published reconstructions of African American and Afro-Caribbean history [Bryc et al., 2015], Latino population structure [Moreno-Estrada et al., 2013], and South Asian admixture chronology [Moorjani et al., 2013], providing independent biological validation for the inferred graph.

6 Conclusion

We have presented HapGraph, an open-source Python tool for joint admixture graph inference and timing estimation from phased SNP data. By combining an F-statistics-based greedy topology search with a Bayesian IBD-based timing model, HapGraph addresses a gap that has persisted in population genetic inference: the lack of a unified framework for estimating who mixed with whom and when, within a single calibrated analysis.

Simulation experiments across three scenarios demonstrate that HapGraph achieves zero false positive rate on pure trees, 100% recall of true admixture target populations, and well-calibrated posterior timing estimates (85–93% empirical coverage for T across $K = 1$ and $K = 2$ scenarios). The key limitation identified in benchmarking — under-calibrated uncertainty for admixture proportions (55–62% empirical coverage) — is a known challenge of the F_3 method-of-moments delta-method approximation near $\alpha \approx 0.5$, and is a priority for future development.

Application to the 1000 Genomes Project 2022 dataset yielded eight biologically interpretable admixture events consistent with published population histories, demonstrating that HapGraph's topology inference scales to 26 populations and produces results suitable for downstream biological interpretation.

Future directions include: extending simulation validation to $K \geq 3$; applying IBD-based timing estimation to real data following hap-ibd pre-processing; improving alpha credible interval calibration; and direct quantitative comparison to GLOBETROTTER and ALDER timing estimates. HapGraph is designed as a modular Python package, allowing components (F-statistics, IBD model, MCMC sampler) to be extended or replaced independently, facilitating community development and adaptation to new use cases.

References

1000 Genomes Project Consortium. A global reference for human genetic variation. *Nature*, 526:68–74, 2015. doi: 10.1038/nature15393.

Oriol Abril-Pla, Virgile Andreani, Colin Carroll, Larry Dong, Christopher J Fonnesbeck, et al. PyMC: a modern and comprehensive probabilistic programming framework in Python. *PeerJ Computer Science*, 9: e1516, 2023. doi: 10.7717/peerj-cs.1516.

Brian L Browning, Sharon R Browning, and Ying Zhou. A fast and simple method for detecting identity-by-descent segments in large-scale data. *American Journal of Human Genetics*, 106:426–437, 2020. doi: 10.1016/j.ajhg.2020.02.010.

Katarzyna Bryc, Eric Y Durand, J Michael Macpherson, David Reich, and Joanna L Mountain. The genetic ancestry of African Americans, Latinos, and European Americans across the United States. *American Journal of Human Genetics*, 96:37–53, 2015. doi: 10.1016/j.ajhg.2014.11.010.

Marta Byrska-Bishop, Uday S Evani, Xuefang Zhao, et al. High-coverage whole-genome sequencing of the expanded 1000 Genomes Project cohort including 602 trios. *Cell*, 185:3426–3440, 2022. doi: 10.1016/j.cell.2022.08.004.

Wolfgang Haak, Iosif Lazaridis, Nick Patterson, et al. Massive migration from the steppe was a source for Indo-European languages in Europe. *Nature*, 522:207–211, 2015. doi: 10.1038/nature14317.

Garrett Hellenthal, George B J Busby, Gavin Band, James F Wilson, Cristian Capelli, Daniel Falush, and Simon Myers. A genetic atlas of human admixture history. *Science*, 343:747–751, 2014. doi: 10.1126/science.1243518.

Matthew D Hoffman and Andrew Gelman. The No-U-Turn sampler: adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15:1593–1623, 2014.

Jerome Kelleher, Alison M Etheridge, and Gilean McVean. Efficient coalescent simulation and genealogical analysis for large sample sizes. *PLoS Computational Biology*, 12:e1004842, 2016. doi: 10.1371/journal.pcbi.1004842.

Po-Ru Loh, Mark Lipson, Nick Patterson, Priya Moorjani, Joseph K Pickrell, David Reich, and Bonnie Berger. Inferring admixture histories of human populations using linkage disequilibrium. *Genetics*, 193:1233–1254, 2013. doi: 10.1534/genetics.112.147330.

Robert Maier, Pavel Flegontov, Olga Flegontova, Ulas Isildak, Piya Changmai, and David Reich. On the limits of fitting complex models of population history to F-statistics. *eLife*, 12:e85492, 2023. doi: 10.7554/eLife.85492.

Alistair Miles, Peter Ralph, Sorin Rae, and Rahul Pisupati. cggh/scikit-allel: v1.3.5. *Zenodo*, 2021. doi: 10.5281/zenodo.3935371.

Priya Moorjani, Kumarasamy Thangaraj, Nick Patterson, Mark Lipson, Po-Ru Loh, Periyasamy Govindaraj, Bonnie Berger, David Reich, and Lalji Singh. Genetic evidence for recent population mixture in India. *American Journal of Human Genetics*, 93:422–438, 2013. doi: 10.1016/j.ajhg.2013.07.006.

Andres Moreno-Estrada, Simon Gravel, Fouad Zakharia, Jacob L McCauley, Jake K Byrnes, Christopher R Gignoux, Patricia A Ortiz-Tello, Ricardo J Martinez, Dale J Hedges, et al. Reconstructing the population genetic history of the Caribbean. *PLoS Genetics*, 9:e1003925, 2013. doi: 10.1371/journal.pgen.1003925.

Svend V Nielsen, Andrew H Vaughn, Kaisa Leppala, Michael J Landis, Thomas Mailund, and Rasmus Nielsen. Bayesian inference of admixture graphs on Native American and Arctic populations. *PLoS Genetics*, 19: e1010410, 2023. doi: 10.1371/journal.pgen.1010410.

Pier Francesco Palamara, Todd Lencz, Ariel Darvasi, and Itsik Pe'er. Length distributions of identity by descent reveal fine-scale demographic history. *American Journal of Human Genetics*, 91:809–822, 2012. doi: 10.1016/j.ajhg.2012.08.030.

Nick Patterson, Priya Moorjani, Yong Luo, Swapan Mallick, Nadin Rohland, Yiping Zhan, Teri Genschoreck, Teresa Webster, and David Reich. Ancient admixture in human history. *Genetics*, 192:1065–1093, 2012. doi: 10.1534/genetics.112.145037.

Joseph K Pickrell and Jonathan K Pritchard. Inference of population splits and mixtures from genome-wide allele frequency data. *PLoS Genetics*, 8:e1002967, 2012. doi: 10.1371/journal.pgen.1002967.

David Reich, Kumarasamy Thangaraj, Nick Patterson, Alkes L Price, and Lalji Singh. Reconstructing Indian population history. *Nature*, 461:489–494, 2009. doi: 10.1038/nature08365.

Naruya Saitou and Masatoshi Nei. The neighbor-joining method: a new method for reconstructing phylogenetic trees. *Molecular Biology and Evolution*, 4:406–425, 1987. doi: 10.1093/oxfordjournals.molbev.a040454.